

Introduction to NoSQL

Nicolas Travers
CNAM – France

Schedule & Organization



- Introduction to **NoSQL** databases
 - 3V, ACID vs BASE, families, CAP theorem, JSon
- Presentation of **MongoDB**
 - Language, distribution, replication, application
- **Practice Works** on MongoDB
 - Queries : find + aggregate
- Material available at: <http://www.chewbii.com/ESSEC>

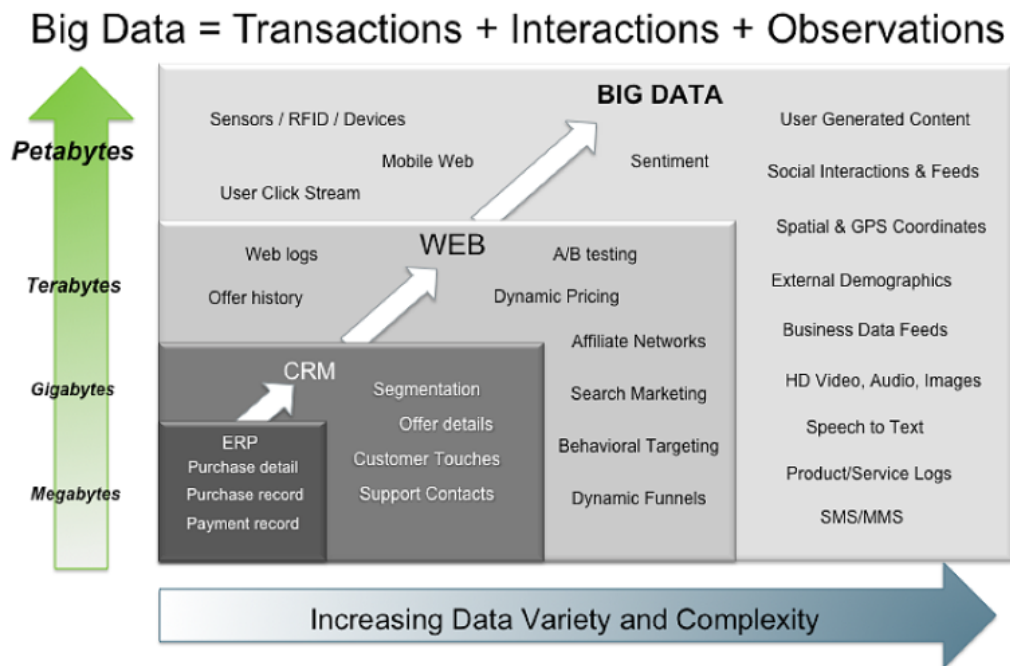
DBMS vs NoSQL

Fault-tolerance

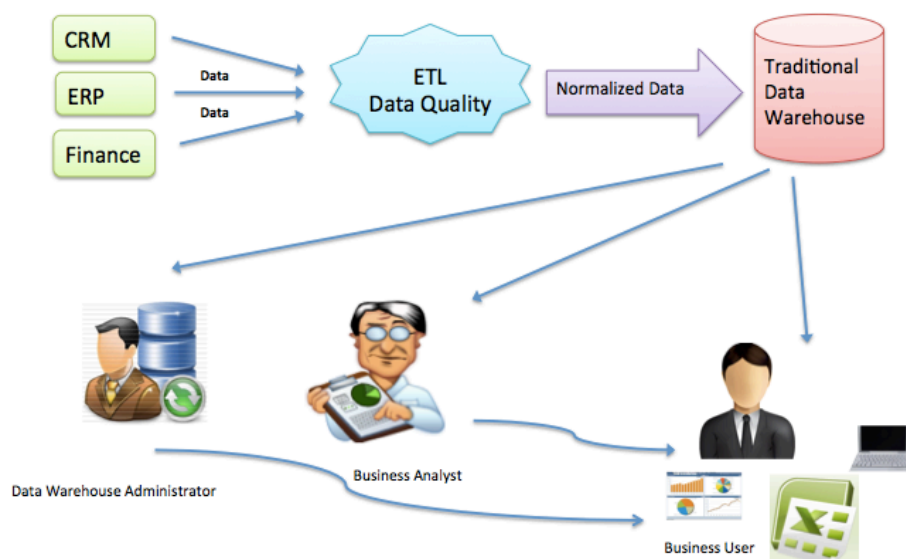


Context

- Applications and Web platforms
 - Exponential growth of the amount of Data (x2 / 2 years)
 - Unprecedented management of this volume
 - Need to distribute both computation and data
 - Huge number of servers
 - Heterogeneous data, maybe complex and often linked
- Ex :
 - Google, Amazon, Facebook
 - Google DataCenter :
 - 5000 servers/data center, ~1M de servers
 - Facebook :
 - 1 PetaBytes of data



BI : Traditional methods



Decisional vs 3V

Incompatible classical approach with the **3V** of *BigData* :

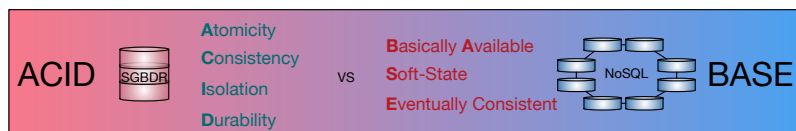
- **Volume**: Designed to store GB/TB of data, but needs PB (maybe EB).
- **Variety**: Heterogeneous and variable types of data, text, semi-structured
- **Velocity**: Data are produced more and more quickly

DBMS: Limitations

- Standard databases
 - Functionalities
 - Joins between tables
 - Complex queries
 - Strong coherency of data
- Requirements in a distributed context
 - Links between entities => same server
 - ++ links => difficulties for data organization

RDBMS vs Distribution

- ACID vs BASE
- **ACID** properties for transactions
 - **Atomicity**: integral completion or none
 - **Consistency**: consistent at start and end
 - **Isolation**: no communication between them
 - **Durability**: an operation cannot be reversed
- Systèmes distribués : modèle **BASE**
 - **Basically Available** :
 - Any request => An answer
 - Even in a changing state
 - **Soft-state** :
 - Opposite to Durability.
 - System's state (servers or data) could change over time (without any update)
 - **Eventually consistent** :
 - With time, data can be consistent
 - Updates have to be propagated




Solution: NoSQL

- NoSQL : **Not Only SQL**
 - New data storage/management approach
 - Scales up the system (through distribution)
 - Complex metadata management
 - No schema
- Do not substitute DBMS, dedicated to:
 - Very huge volume of data (PetaBytes)
 - Very short response time
 - Consistency is not mandatory

Databases and NoSQL

The Hard Life of a NoSQL Coder



Part 1: The Outing

NoSQL DB: Characteristics


- **No relations** => *Collections*
 - No fix structures (nay none)
- **Complex data** (e.g. *documents*)
 - Objects, nesting, arrays
- **Data distribution**
 - High parallelization (Map/Reduce)
- **Data replication**
 - Disponibility vs Consistency (no transactions)
 - Few writes, many reads

Sharding : Scalability

- Datablocks are distributed in a cluster of servers
- Horizontal partitioning
- 3 types of technics:
 1. Resource allocation based: *HDFS*
 - ++ Fault tolerance and massive computations
 2. Tree-based structure: *Clustered index* (sort)
 - ++ Physical data clustering and dynamicity
 3. Hash-based structure: *Consistent Hashing*
 - ++ elasticity and self-management

Several NoSQL systems

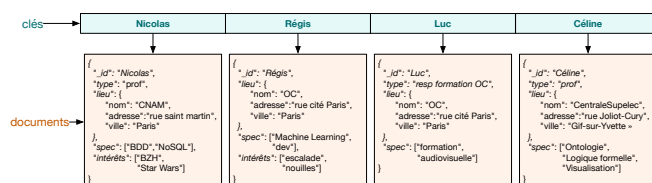
Key-Value stores




Column oriented

id	type	id	lieu	id	spec	id	intérêts
Nicolas	prof	Céline	Centrale Supélec	Nicolas	BDD	Nicolas	BZH
Céline	prof	Nicolas	CNAM	Nicolas	NoSQL	Nicolas	Star Wars
Luc	resp formation OC	Régis	OC	Régis	Machine Learning	Régis	escalade
		Luc	OC	Régis	Dev	Régis	nouilles chinoises
				Luc	formation		
				Luc	audiovisuel		
				Céline	Ontologie		
				Céline	logique formelle		
				Céline	visualisation		

Document Oriented



Graph oriented



I - NoSQL & Key-Value store


- Similar to a distributed “*HashMap*”
- Key + Value
 - No fixed schema on values (strings, object, integer, binaries...)
- Drawbacks:
 - No structures nor typing
 - No structured-based queries
- DynamoDB (Amazon), Redis (VMWare), Voldemort (LinkedIn)

I - NoSQL & Key-Value store

Example

Relationnel

id	type	lieu	spec	intérêts
Nicolas	prof	CNAM	BDD, NoSQL	BZH, Star Wars
Régis		OC	Machine Learning, Dev	escalade, nouilles chinoises
Luc	resp formation OC	OC	formation, audiovisuel	
Céline	prof	CentraleSupélec	Ontologie, logique formelle, visualisation	




I - NoSQL & Key-Value store

Queries


• CRUD

- CREATE (**clé**, **valeur**)
 - CREATE ("Nicolas", "type:'prof',lieu:'CNAM',spec:'BDD,NoSQL',interets:'BZH,Star Wars' ") → OK
- READ(**clé**)
 - READ("Nicolas") → "type:'prof',lieu:'CNAM',spec:'BDD,NoSQL',interets:'BZH,Star Wars' "
- UPDATE(**clé**, **valeur**)
 - UPDATE("Nicolas", "type:'prof',lieu:'CNAM,CS',spec:'BDD,NoSQL' ") → OK
- DELETE(**clé**)
 - DELETE("Nicolas") → OK



I - NoSQL & Key-Value store

solutions



Efficiency

Easy to set up

II – NoSQL & Columns

- Column-based storage
 - DBMS: tuples (lines)
 - Easy to insert a new column
 - Dynamic schema
- BigTable/Hbase (Google), Cassandra (Facebook&Apache), SimpleDB (Amazon)

II - NoSQL & Columns

example

id	type	lieu	spec	intérêts
Nicolas	prof	CNAM	BDD, NoSQL	BZH, Star Wars
Régis		OC	Machine Learning, Dev	escalade, nouilles chinoises
Luc	resp formation OC	OC	formation, audiovisuel	
Céline	prof	CentraleSupelec	Ontologie, logique formelle, visualisation	

id	type
Nicolas	prof
Céline	prof
Luc	resp formation OC

id	lieu
Céline	Centrale Supelec
Nicolas	CNAM
Régis	OC
Luc	OC

id	spec
Nicolas	BDD
Nicolas	NoSQL
Régis	Machine Learning
Régis	Dev
Luc	formation
Luc	audiovisuel
Céline	Ontologie
Céline	logique formelle
Céline	visualisation

id	intérêts
Nicolas	BZH
Nicolas	Star Wars
Régis	escalade
Régis	nouilles chinoises

II - NoSQL & Columns

Queries

- Column-oriented queries
 - How many **professors (type)** are at **CentraleSupelec (lieu)**

id	type
Nicolas	prof
Céline	prof

id	lieu
Céline	Centrale Supelec

id	spec
Nicolas	BDD
Nicolas	NoSQL
Régis	Machine Learning
Régis	Dev
Luc	formation
Luc	audiovisuel
Céline	Ontologie
Céline	logique formelle
Céline	visualisation

id	intérêts
Nicolas	BZH
Nicolas	Star Wars
Régis	escalade
Régis	nouilles chinoises

II - NoSQL & Columns

solutions



APACHE
HBASE



Spark SQL

Aggregates

Correlations

III – NoSQL & Documents

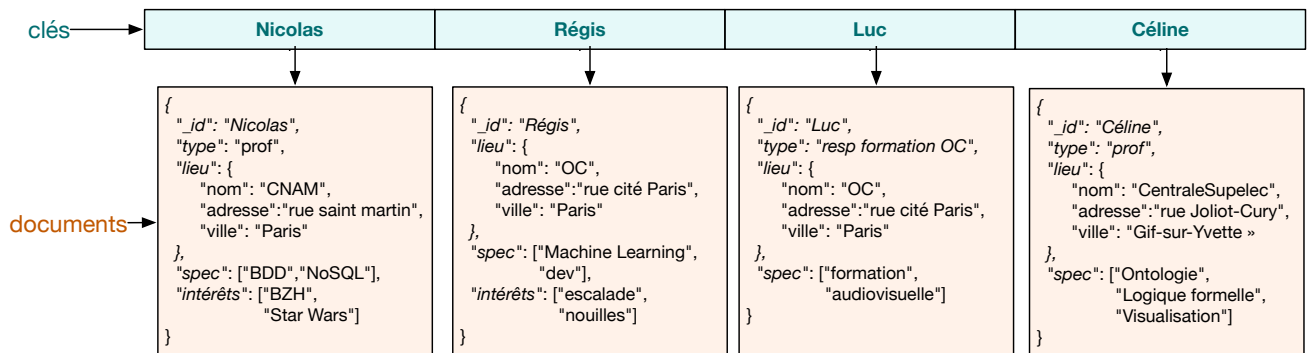
- Based on the key-value store
 - Add semi-structured data (JSON/XML)
- HTTP API
 - More complex than CRUD

➤ MongoDB, CouchDB (Apache), RavenDB, Terrastore

III - NoSQL & Documents

example

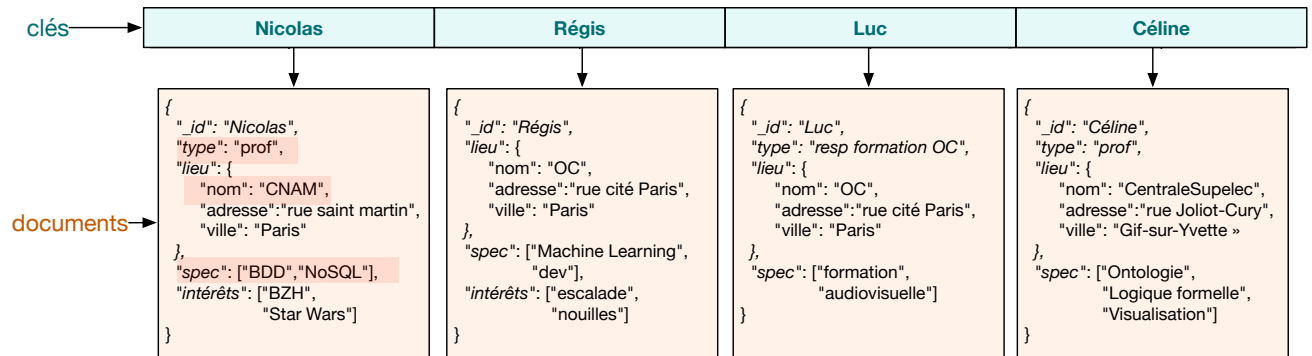
id	type	lieu	spec	intérêts
Nicolas	prof	CNAM	BDD, NoSQL	BZH, Star Wars
Régis		OC	Machine Learning, Dev	escalade, nouilles chinoises
Luc	resp formation OC	OC	formation, audiovisuel	
Céline	prof	CentraleSupélec	Ontologie, logique formelle, visualisation	



III - NoSQL & Documents

Queries

- Manipulations on documents content
 - Establishment (**lieu.nom**) of professors (**type**) specialized in DB (**in spec**)



III - NoSQL & Documents

solutions



Richness of queries

Manage objects

IV – NoSQL & Graph

- Storage: nodes, relations and properties

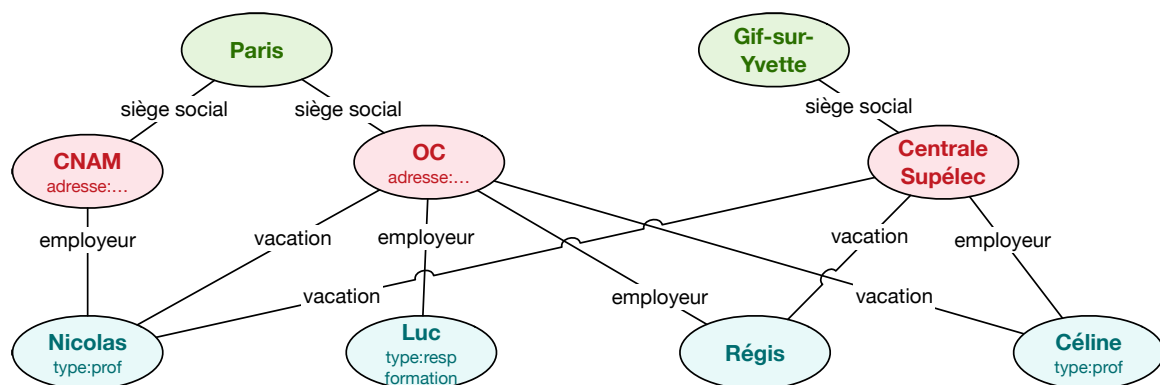
- Graph Theory
- Path querying on the graph
- Data are loaded on demand
- Difficulties for modeling

➤ Neo4j, OrientDB (Apache), FlockDB (Twitter)

IV - NoSQL & Graph

example

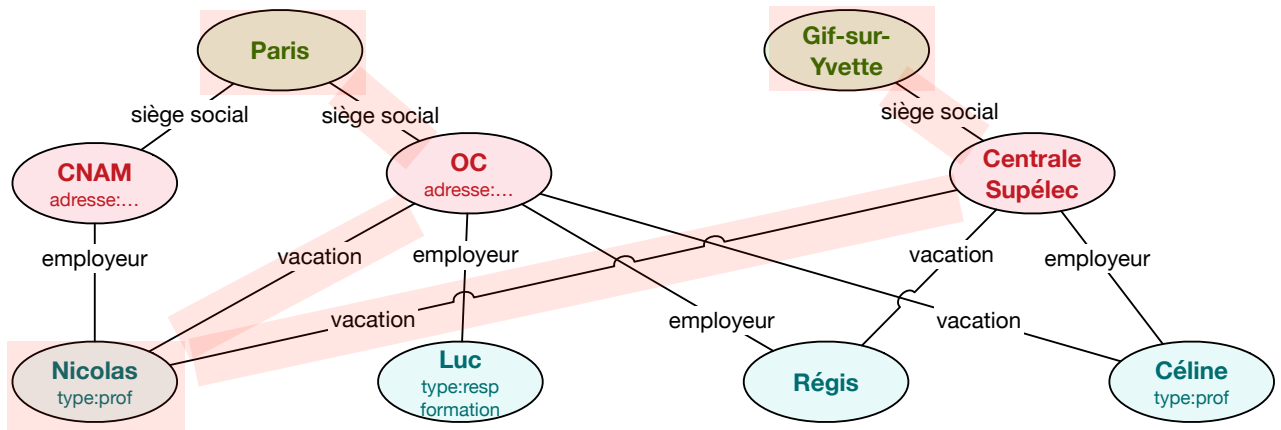
id	type	lieu	spec	intérêts
Nicolas	prof	CNAM	BDD, NoSQL	BZH, Star Wars
Régis		OC	Machine Learning, Dev	escalade, nouilles chinoises
Luc	resp formation OC	OC	formation, audiovisuel	
Céline	prof	CentraleSupélec	Ontologie, logique formelle, visualisation	



IV - NoSQL & Graph

queries

- Pattern queries
 - **Persons** giving **vacations** at **Paris** and **Gif-sur-Yvette**



IV - NoSQL & Graph


solutions




Network



Azure Cosmos DB:

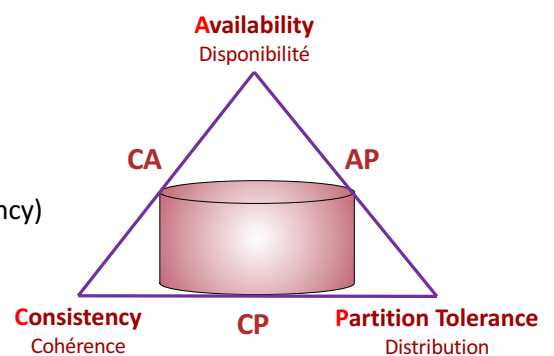


Recommandation



Brewer's CAP Theorem (2000)

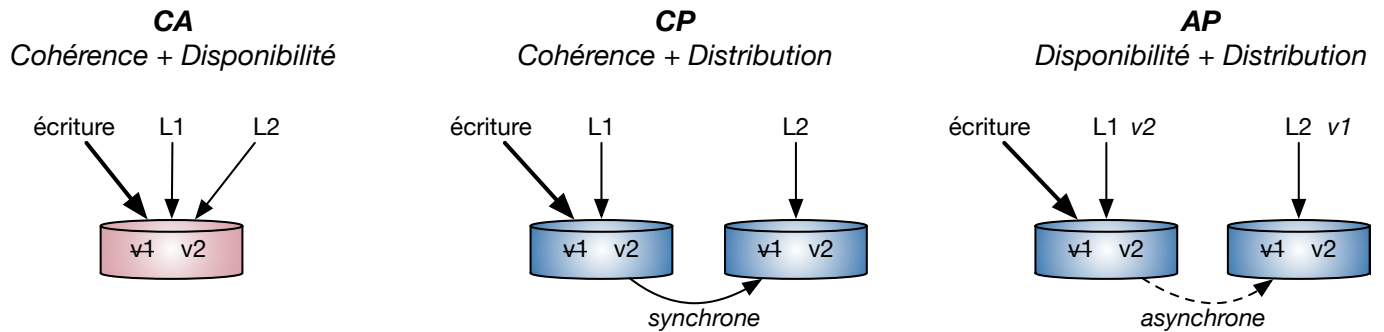
- 3 main properties for distributed management
 - Consistency:**
 - A data have the same value at the same time (coherency)
 - Availability:**
 - Even if a server is down, data is available
 - Partition Tolerance:**
 - Even if the system is partitioned, a query must have an answer (unless for global failures)



- Theorem: A distributed, networked system can have only two of these three properties.

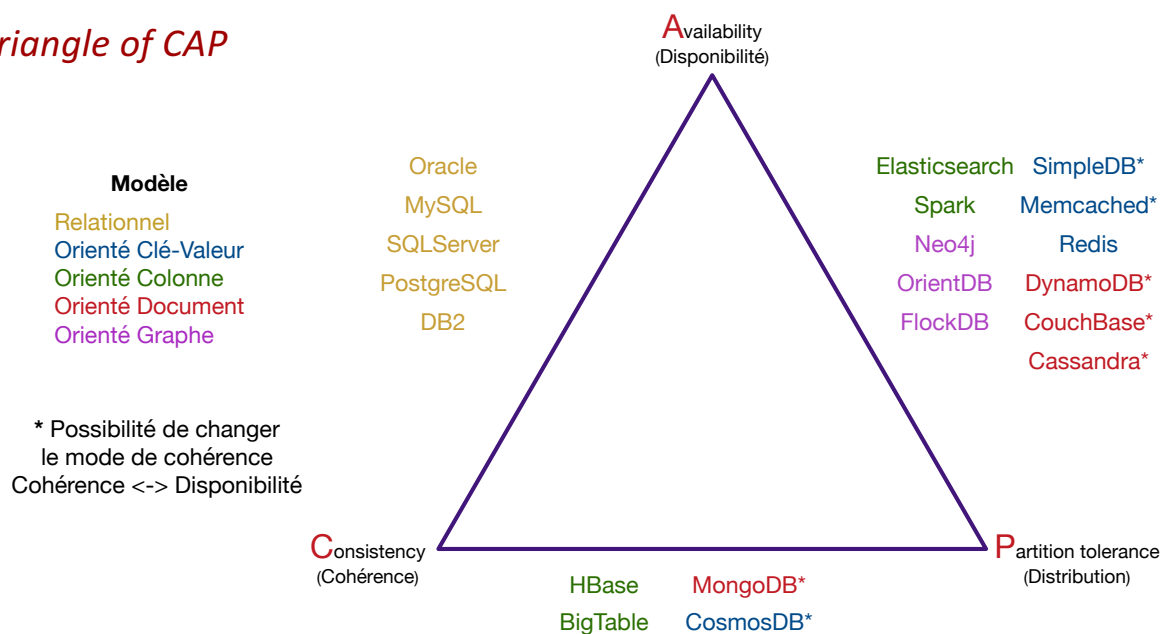
"CAP Theorem"


illustration



"CAP Theorem"

Triangle of CAP





- Initially XML used for complex internet communications (Web Services)
 - Too verbose
- **JSON (JavaScript Object Notation)**
 - Lightweight, text-oriented, language independent
 - Used for several Web services (Google API, Twitter API)

JSON : Structures

- **Key + Value**
 - `"lastname" : "Travers"`
 - Keys with quotations
- **Objects/documents**
 - Collection of key/values
 - `{ "lastname" : "Travers",
 "firstname" : "Nicolas",
 "kind" : 1 }`

Data types

- **Scalar** : String, Integer, float, boolean, null...
- **List** : arrays [...]
- **Documents** : objects {...}

Arrays

- No typing inside arrays
 - “lessons” : [“SQL”, 1, 4.2, null, “NoSQL”]
- Can nest documents
 - “doc” : [{“test” : 1},
 {“test” : {“nesting” : 1.0}},
 {“key” : “text”, “value” : null}]

JSON : Identifiers

- Key « `_id` » commonly used to identify documents
 - Overwrite already stored ids
 - Can be automatically generated
 - Ex MongoDB : “_id” : ObjectId(1234567890)

Json : complete example

```
{
  "_id" : 1234,
  "lastname" : "Travers", "firstname" : "Nicolas",
  "work" : {
    "company" : "Cnam",
    "location" : {
      "street" : "2 rue cont  ",
      "city" : "Paris",
      "zip" : 75141
    },
  },
  "fields" : [ "DB", "DB tuning", "XML", "NoSQL", "IR" ]
}
```